

Finding the MRI brain tumor detection dataset

Let's find the dataset in this link: <https://www.kaggle.com/navoneel/brain-mri-images-for-brain-tumor-detection> (<https://www.kaggle.com/navoneel/brain-mri-images-for-brain-tumor-detection>)

Import packages

```
In [20]: import numpy as np
import torch
from torch.utils.data import Dataset, DataLoader, ConcatDataset
import glob2
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, accuracy_score
import random
import cv2
import sys
```

Reading the Images

```
In [21]: tumor = []
healthy = []
for f in glob2.iglob("./data/brain_tumor_dataset/yes/*.jpg"):
    img = cv2.imread(f)
    img = cv2.resize(img,(128,128))
    b, g, r = cv2.split(img)
    img = cv2.merge([r,g,b])
    tumor.append(img)

for f in glob2.iglob("./data/brain_tumor_dataset/no/*.jpg"):
    img = cv2.imread(f)
    img = cv2.resize(img,(128,128))
    b, g, r = cv2.split(img)
    img = cv2.merge([r,g,b])
    healthy.append(img)
```

```
In [4]: healthy = np.array(healthy)
tumor = np.array(tumor)
```

```
In [5]: healthy.shape
```

```
Out[5]: (91, 128, 128, 3)
```

```
In [6]: tumor.shape
```

```
Out[6]: (154, 128, 128, 3)
```

```
In [7]: np.random.choice(10, 5, replace=False)
```

```
Out[7]: array([5, 6, 7, 9, 8])
```

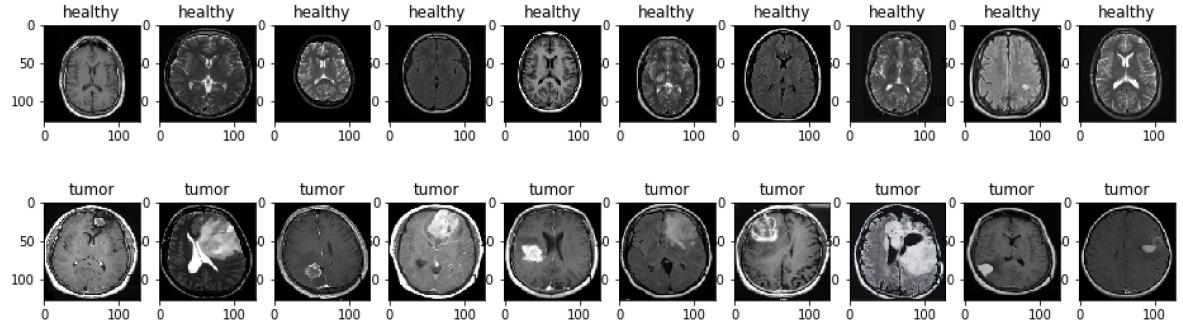
Visualizing Brain MRI Images

```
In [8]: def plot_random(healthy, tumor, num=5):
    healthy_imgs = healthy[np.random.choice(healthy.shape[0], num, replace=False)]
    tumor_imgs = tumor[np.random.choice(tumor.shape[0], num, replace=False)]

    plt.figure(figsize=(16,9))
    for i in range(num):
        plt.subplot(1, num, i+1)
        plt.title('healthy')
        plt.imshow(healthy_imgs[i])

    plt.figure(figsize=(16,9))
    for i in range(num):
        plt.subplot(1, num, i+1)
        plt.title('tumor')
        plt.imshow(tumor_imgs[i])
```

```
In [9]: plot_random(healthy, tumor, num=10)
```



Create Torch Dataset Class

What is Pytorch's Abstract Dataset Class

```
In [10]: class Dataset(object):
    """An abstract class representing a Dataset.

    All other datasets should subclass it. All subclasses should override
    ``__len__``, that provides the size of the dataset, and ``__getitem__``,
    supporting integer indexing in range from 0 to len(self) exclusive.
    """

    def __getitem__(self, index):
        raise NotImplementedError

    def __len__(self):
        raise NotImplementedError

    def __add__(self, other):
        return ConcatDataset([self, other])
```

Creating MRI custom dataset class

```
In [22]: class MRI(Dataset):
    def __init__(self):

        tumor = []
        healthy = []
        # cv2 - It reads in BGR format by default
        for f in glob2.iglob("./data/brain_tumor_dataset/yes/*.jpg"):
            img = cv2.imread(f)
            img = cv2.resize(img,(128,128)) # I can add this later in the boot
            b, g, r = cv2.split(img)
            img = cv2.merge([r,g,b])
            img = img.reshape((img.shape[2],img.shape[0],img.shape[1])) # other
            tumor.append(img)

        for f in glob2.iglob("./data/brain_tumor_dataset/no/*.jpg"):
            img = cv2.imread(f)
            img = cv2.resize(img,(128,128))
            b, g, r = cv2.split(img)
            img = cv2.merge([r,g,b])
            img = img.reshape((img.shape[2],img.shape[0],img.shape[1]))
            healthy.append(img)

        # our images
        tumor = np.array(tumor,dtype=np.float32)
        healthy = np.array(healthy,dtype=np.float32)

        # our labels
        tumor_label = np.ones(tumor.shape[0], dtype=np.float32)
        healthy_label = np.zeros(healthy.shape[0], dtype=np.float32)

        # Concatenates
        self.images = np.concatenate((tumor, healthy), axis=0)
        self.labels = np.concatenate((tumor_label, healthy_label))

    def __len__(self):
        return self.images.shape[0]

    def __getitem__(self, index):

        sample = {'image': self.images[index], 'label':self.labels[index]}

        return sample

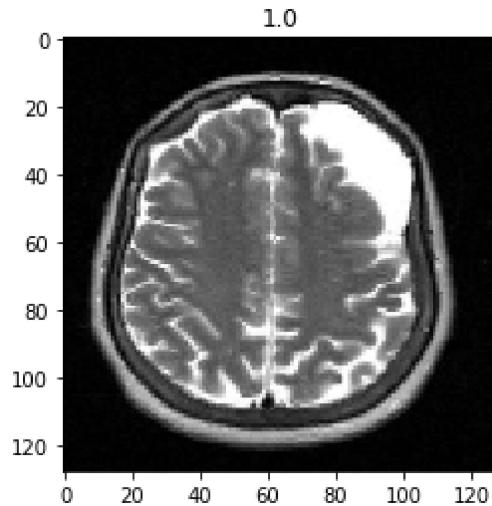
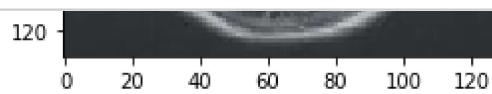
    def normalize(self):
        self.images = self.images/255.0
```

```
In [17]: mri_dataset = MRI()
mri_dataset.normalize()
```

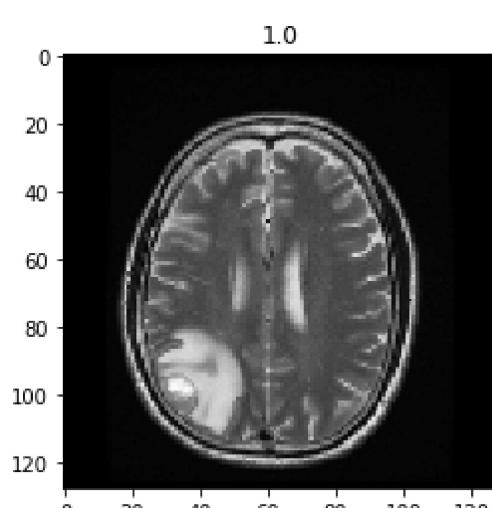
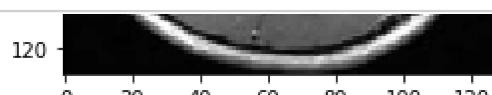
Creating a dataloader

The World Without dataloaders

```
In [23]: index = list(range(len(mri)))
random.shuffle(index)
for idx in index:
    sample = mri[idx]
    img = sample['image']
    label = sample['label']
    img = img.reshape(img.shape[1], img.shape[2], img.shape[0])
    plt.title(label)
    plt.imshow(img)
    plt.show()
```



```
In [29]: it = iter(mri)
for i in range(10):
    sample = next(it)
    img = sample['image']
    label = sample['label']
    img = img.reshape(img.shape[1], img.shape[2], img.shape[0])
    plt.title(label)
    plt.imshow(img)
    plt.show()
```



The World with data loaders

```
In [30]: dataloader = DataLoader (mri_dataset, batch_size=10, shuffle=True)
```

```
In [32]: for sample in dataloader:  
    img = sample['image']  
    print(img.shape)  
    sys.exit()  
    #img = img.reshape(img.shape[1], img.shape[2], img.shape[0])  
    #plt.imshow(img)  
    #plt.show()
```

```
torch.Size([10, 3, 128, 128])
```

```
An exception has occurred, use %tb to see the full traceback.
```

```
SystemExit
```

```
C:\Users\srrira\anaconda3\envs\MRI\lib\site-packages\IPython\core\interactiveshell.py:3304: UserWarning: To exit: use 'exit', 'quit', or Ctrl-D.  
    warn("To exit: use 'exit', 'quit', or Ctrl-D.", stacklevel=1)
```

Create a model

```
In [33]: import torch.nn as nn
import torch.nn.functional as F

class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.cnn_model = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=6, kernel_size=5),
            nn.Tanh(),
            nn.AvgPool2d(kernel_size=2, stride=5),
            nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5),
            nn.Tanh(),
            nn.AvgPool2d(kernel_size=2, stride=5))

        self.fc_model = nn.Sequential(
            nn.Linear(in_features=256, out_features=120),
            nn.Tanh(),
            nn.Linear(in_features=120, out_features=84),
            nn.Tanh(),
            nn.Linear(in_features=84, out_features=1))

    def forward(self, x):
        x = self.cnn_model(x)
        x = x.view(x.size(0), -1)
        x = self.fc_model(x)
        x = F.sigmoid(x)

    return x
```

The logic behind the numbers

Here is the grand formula

$$n_{out} = \text{ceil}\left[\frac{n_{in} + 2p - f}{s} + 1\right]$$

Look into the parameters of the model

```
In [34]: model = CNN()
```

In [35]: model

```
Out[35]: CNN(  
    (cnn_model): Sequential(  
        (0): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))  
        (1): Tanh()  
        (2): AvgPool2d(kernel_size=2, stride=5, padding=0)  
        (3): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))  
        (4): Tanh()  
        (5): AvgPool2d(kernel_size=2, stride=5, padding=0)  
    )  
    (fc_model): Sequential(  
        (0): Linear(in_features=256, out_features=120, bias=True)  
        (1): Tanh()  
        (2): Linear(in_features=120, out_features=84, bias=True)  
        (3): Tanh()  
        (4): Linear(in_features=84, out_features=1, bias=True)  
    )  
)
```

In [38]: model.cnn_model

```
Out[38]: Sequential(  
    (0): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))  
    (1): Tanh()  
    (2): AvgPool2d(kernel_size=2, stride=5, padding=0)  
    (3): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))  
    (4): Tanh()  
    (5): AvgPool2d(kernel_size=2, stride=5, padding=0)  
)
```

In [39]:

```
model.cnn_model[0]
```

```
Out[39]: Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
```

In [40]: `model.cnn_model[0].weight`

```
[[ 2.8615e-02,  1.4409e-04, -1.9788e-03,  4.6759e-02,  7.1333e-0
2],
 [-1.1318e-01, -6.9383e-02,  5.9661e-02,  5.0969e-02,  9.3040e-0
2],
 [-1.0679e-01,  1.0443e-01,  1.0835e-02, -2.4515e-02, -5.2875e-0
3],
 [-1.0584e-01, -1.2717e-02, -8.8371e-02, -5.1271e-02, -7.8316e-0
2],
 [-1.0457e-01, -7.9666e-02,  9.5497e-02, -1.5778e-02, -6.5364e-0
2]],

[[ 8.0895e-02,  8.7384e-02,  7.9559e-02, -3.6089e-02, -9.6216e-0
2],
 [ 8.4533e-02,  2.3745e-03, -1.0332e-01, -1.1450e-01,  1.0005e-0
1],
 [-4.6434e-02,  7.2655e-02, -1.1040e-01,  1.3803e-02,  1.3546e-0
2],
 [ 9.8149e-02, -1.1156e-01,  5.7886e-03,  6.9954e-02, -8.1929e-0
3],
 [-0.2071e-02,  0.0121e-02, -1.0060e-02, -2.2216e-02, -0.7811e-0
2],
```

In [42]: `model.cnn_model[0].weight[0].shape`

Out[42]: `torch.Size([3, 5, 5])`

In [43]: `model.cnn_model[0].weight[0][1]`

```
tensor([[ 0.0286,  0.0001, -0.0020,  0.0468,  0.0713],
 [-0.1132, -0.0694,  0.0597,  0.0510,  0.0930],
 [-0.1068,  0.1044,  0.0108, -0.0245, -0.0053],
 [-0.1058, -0.0127, -0.0884, -0.0513, -0.0783],
 [-0.1046, -0.0797,  0.0955, -0.0158, -0.0654]], grad_fn=<SelectBackward0>)
```

Linear layer

In [44]: `model.fc_model`

```
Sequential(
(0): Linear(in_features=256, out_features=120, bias=True)
(1): Tanh()
(2): Linear(in_features=120, out_features=84, bias=True)
(3): Tanh()
(4): Linear(in_features=84, out_features=1, bias=True)
)
```

In [45]: `model.fc_model[0]`

Out[45]: `Linear(in_features=256, out_features=120, bias=True)`

In [46]: `model.fc_model[0].weight.shape`

Out[46]: `torch.Size([120, 256])`

Understanding x.view(x.size(0),-1)

```
In [48]: x = torch.tensor ([1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16])
```

```
In [49]: x = x.reshape((2, 2, 2, 2))
```

```
In [51]: x
```

```
Out[51]: tensor([[[[ 1,  2],
                   [ 3,  4]],

                  [[ 5,  6],
                   [ 7,  8]],

                  [[[ 9, 10],
                   [11, 12]],

                  [[13, 14],
                   [15, 16]]]])
```

```
In [52]: x.size(0)
```

```
Out[52]: 2
```

```
In [53]: x.view(-1)
```

```
Out[53]: tensor([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16])
```

```
In [54]: x.view(x.size(0), -1).shape
```

```
Out[54]: torch.Size([2, 8])
```

Some Basics of Training and Evaluation in Pytorch

model.eval()

- Used particularly for inference **NOTHING to DO with gradients!!!**
- changes the forward() behaviour of the module it is called up on eg, it disables dropout and has batch norm use the entire population statistics. This is necessary for inference

model.train()

- Brings drop out and batch norm to action (i.e., train mode).
- Gradients are computed

numpy array vs tensor

The difference between a NumPy array and a tensor is that the tensors are backed by the accelerator memory like GPU and they are immutable, unlike NumPy arrays. You can never update a tensor but create a new one. If you are into machine learning or going to be into it, A Tensor is a suitable choice if you are going to use GPU. A tensor can reside in accelerator's memory.

- The numpy arrays are the core functionality of the numpy package designed to support faster mathematical operations. Unlike python's inbuilt list data structure, they can only hold elements of a single data type. Library like pandas which is used for data preprocessing is built around the numpy array. **Pytorch tensors are similar to numpy arrays, but can also be operated on CUDA-capable Nvidia GPU.**
- Numpy arrays are mainly used in typical machine learning algorithms (such as k-means or Decision Tree in scikit-learn) whereas pytorch tensors are mainly used in deep learning which requires heavy matrix computation.
- Unlike numpy arrays, while creating pytorch tensor, it also accepts two other arguments called the device_type (whether the computation happens on CPU or GPU) and the requires_grad (which is used to compute the derivatives).

torch.tensor vs. torch.cuda.tensor

The key difference is just that torch.Tensor occupies CPU memory while torch.cuda.Tensor occupies GPU memory. Of course operations on a CPU Tensor are computed with CPU while operations for the GPU / CUDA Tensor are computed on GPU.

```
In [55]: # device will be 'cuda' if a GPU is available
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

# creating a CPU tensor
cpu_tensor = torch.rand(10).to(device)
# moving same tensor to GPU
gpu_tensor = cpu_tensor.to(device)

print(cpu_tensor, cpu_tensor.dtype, type(cpu_tensor), cpu_tensor.type())
print(gpu_tensor, gpu_tensor.dtype, type(gpu_tensor), gpu_tensor.type())

print(cpu_tensor*gpu_tensor)

tensor([0.5118, 0.3718, 0.2437, 0.9759, 0.9754, 0.9759, 0.4781, 0.8415, 0.285
0,
       0.6792]) torch.float32 <class 'torch.Tensor'> torch.FloatTensor
tensor([0.5118, 0.3718, 0.2437, 0.9759, 0.9754, 0.9759, 0.4781, 0.8415, 0.285
0,
       0.6792]) torch.float32 <class 'torch.Tensor'> torch.FloatTensor
tensor([0.2619, 0.1383, 0.0594, 0.9523, 0.9514, 0.9524, 0.2286, 0.7081, 0.081
2,
       0.4613])
```

As the underlying hardware interface is completely different, CPU Tensors are just compatible with CPU Tensor and vice versa GPU Tensors are just compatible to GPU Tensors.

In which scenario is torch.cuda.Tensor() necessary?

When you want to use GPU acceleration (which is much faster in most cases) for your program, you need to use torch.cuda.Tensor, but you have to make sure that ALL tensors you are using are CUDA Tensors, mixing is not possible here.

tensor.cpu().detach().numpy(): Convert Pytorch tensor to Numpy array

As mentioned before, np.ndarray object does not have this extra "computational graph" layer and therefore, when converting a torch.tensor to np.ndarray you must explicitly remove the computational graph of the tensor using the detach() command. .cpu() returns a copy of this object in CPU memory.

Evaluate a New-Born Neural Network!

```
In [59]: mri_dataset = MRI()  
mri_dataset.normalize()  
device = torch.device('cpu')  
model = CNN().to(device)
```

```
In [60]: dataloader = DataLoader(mri_dataset, batch_size=32, shuffle=False)
```

```
In [61]: model.eval()
outputs = []
y_true = []
with torch.no_grad():
    for D in dataloader:
        image = D['image'].to(device)
        label = D['label'].to(device)

        y_hat = model(image)

        outputs.append(y_hat.cpu().detach().numpy())
        y_true.append(label.cpu().detach().numpy())
```

```
C:\Users\srira\anaconda3\envs\MRI\lib\site-packages\torch\nn\functional.py:18
06: UserWarning: nn.functional.sigmoid is deprecated. Use torch.sigmoid instead.
    warnings.warn("nn.functional.sigmoid is deprecated. Use torch.sigmoid instead.")
```

```
In [62]: outputs = np.concatenate(outputs, axis=0).squeeze()
y_true = np.concatenate(y_true, axis=0).squeeze()
```

```
In [64]: def threshold(scores, threshold=0.50, minimum=0, maximum = 1.0):
    x = np.array(list(scores))
    x[x >= threshold] = maximum
    x[x < threshold] = minimum
    return x
```

```
In [65]: accuracy_score(y_true, threshold(outputs))
```

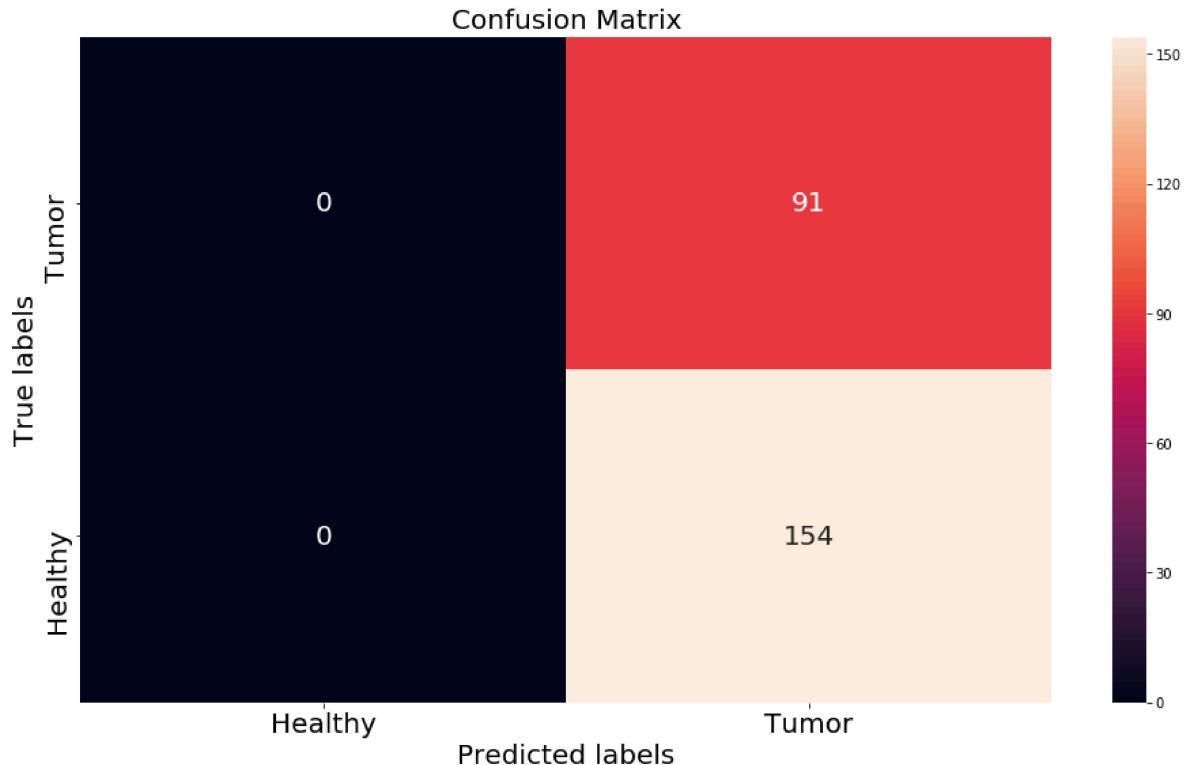
```
Out[65]: 0.6285714285714286
```

```
In [66]: # a better confusion matrix
import seaborn as sns

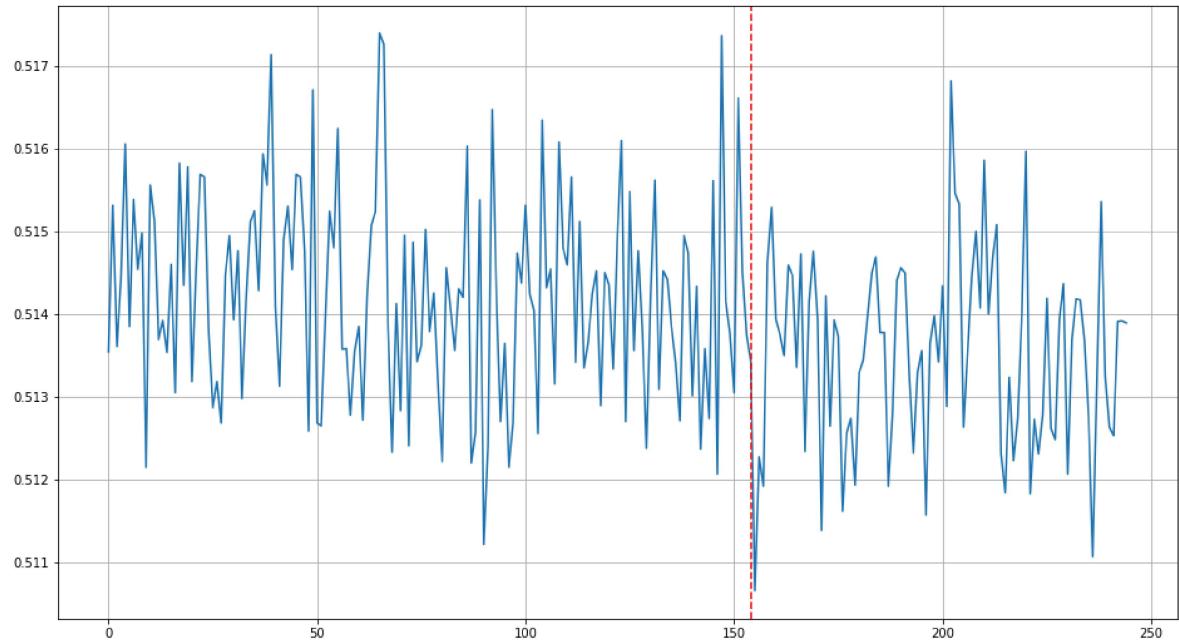
plt.figure(figsize=(16,9))
cm = confusion_matrix(y_true, threshold(outputs))
ax= plt.subplot()
sns.heatmap(cm, annot=True, fmt='g', ax=ax, annot_kws={"size": 20})

# labels, title and ticks
ax.set_xlabel('Predicted labels', fontsize=20)
ax.set_ylabel('True labels', fontsize=20)
ax.set_title('Confusion Matrix', fontsize=20)
ax.xaxis.set_ticklabels(['Healthy', 'Tumor'], fontsize=20)
ax.yaxis.set_ticklabels(['Tumor', 'Healthy'], fontsize=20)
```

Out[66]: [Text(0, 0.5, 'Tumor'), Text(0, 1.5, 'Healthy')]



```
In [67]: plt.figure(figsize=(16,9))
plt.plot(outputs)
plt.axvline(x=len(tumor), color='r', linestyle='--')
plt.grid()
```



Train the dumb model

```
In [68]: eta = 0.0001
EPOCH = 400
optimizer = torch.optim.Adam(model.parameters(), lr=eta)
dataloader = DataLoader(mri_dataset, batch_size=32, shuffle=True)
model.train()
```

```
Out[68]: CNN(
  (cnn_model): Sequential(
    (0): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
    (1): Tanh()
    (2): AvgPool2d(kernel_size=2, stride=5, padding=0)
    (3): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
    (4): Tanh()
    (5): AvgPool2d(kernel_size=2, stride=5, padding=0)
  )
  (fc_model): Sequential(
    (0): Linear(in_features=256, out_features=120, bias=True)
    (1): Tanh()
    (2): Linear(in_features=120, out_features=84, bias=True)
    (3): Tanh()
    (4): Linear(in_features=84, out_features=1, bias=True)
  )
)
```

```
In [69]: for epoch in range(1, EPOCH):
    losses = []
    for D in dataloader:
        optimizer.zero_grad()
        data = D['image'].to(device)
        label = D['label'].to(device)
        y_hat = model(data)
        # define Loss function
        error = nn.BCELoss()
        loss = torch.sum(error(y_hat.squeeze(), label))
        loss.backward()
        optimizer.step()
        losses.append(loss.item())
    if (epoch+1) % 10 == 0:
        print('Train Epoch: {} \tLoss: {:.6f}'.format(epoch+1, np.mean(losses)))
```

```
C:\Users\srira\anaconda3\envs\MRI\lib\site-packages\torch\nn\functional.py:18
06: UserWarning: nn.functional.sigmoid is deprecated. Use torch.sigmoid instead.
    warnings.warn("nn.functional.sigmoid is deprecated. Use torch.sigmoid instead.")
```

```
Train Epoch: 10 Loss: 0.613238
Train Epoch: 20 Loss: 0.559900
Train Epoch: 30 Loss: 0.536614
Train Epoch: 40 Loss: 0.522238
Train Epoch: 50 Loss: 0.511623
Train Epoch: 60 Loss: 0.501433
Train Epoch: 70 Loss: 0.484001
Train Epoch: 80 Loss: 0.478290
Train Epoch: 90 Loss: 0.451469
Train Epoch: 100 Loss: 0.430972
Train Epoch: 110 Loss: 0.416857
Train Epoch: 120 Loss: 0.397001
Train Epoch: 130 Loss: 0.379315
Train Epoch: 140 Loss: 0.357959
Train Epoch: 150 Loss: 0.341558
Train Epoch: 160 Loss: 0.308392
Train Epoch: 170 Loss: 0.289783
Train Epoch: 180 Loss: 0.265693
Train Epoch: 190 Loss: 0.243074
Train Epoch: 200 Loss: 0.227914
Train Epoch: 210 Loss: 0.199677
Train Epoch: 220 Loss: 0.189192
Train Epoch: 230 Loss: 0.162135
Train Epoch: 240 Loss: 0.137883
Train Epoch: 250 Loss: 0.120065
Train Epoch: 260 Loss: 0.105881
Train Epoch: 270 Loss: 0.086956
Train Epoch: 280 Loss: 0.071348
Train Epoch: 290 Loss: 0.061062
Train Epoch: 300 Loss: 0.045453
Train Epoch: 310 Loss: 0.037165
Train Epoch: 320 Loss: 0.031246
Train Epoch: 330 Loss: 0.025958
Train Epoch: 340 Loss: 0.021502
Train Epoch: 350 Loss: 0.017400
Train Epoch: 360 Loss: 0.014252
Train Epoch: 370 Loss: 0.012353
Train Epoch: 380 Loss: 0.010776
Train Epoch: 390 Loss: 0.009125
Train Epoch: 400 Loss: 0.007856
```

Evaluate a smart model

```
In [70]: model.eval()
dataloader = DataLoader(mri_dataset, batch_size=32, shuffle=False)
outputs = []
y_true = []
with torch.no_grad():
    for D in dataloader:
        image = D['image'].to(device)
        label = D['label'].to(device)

        y_hat = model(image)

        outputs.append(y_hat.cpu().detach().numpy())
        y_true.append(label.cpu().detach().numpy())

outputs = np.concatenate(outputs, axis=0)
y_true = np.concatenate(y_true, axis=0)
```

```
In [71]: accuracy_score(y_true, threshold(outputs))
```

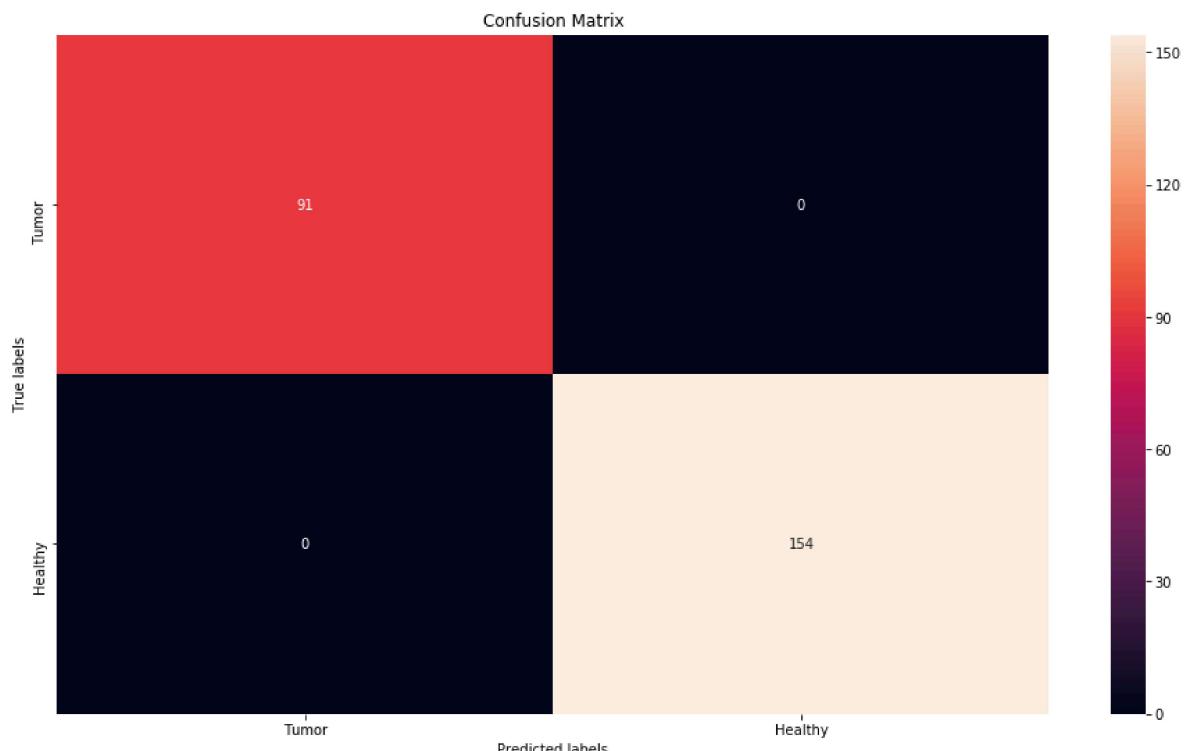
```
Out[71]: 1.0
```

```
In [72]: cm = confusion_matrix(y_true, threshold(outputs))
plt.figure(figsize=(16,9))

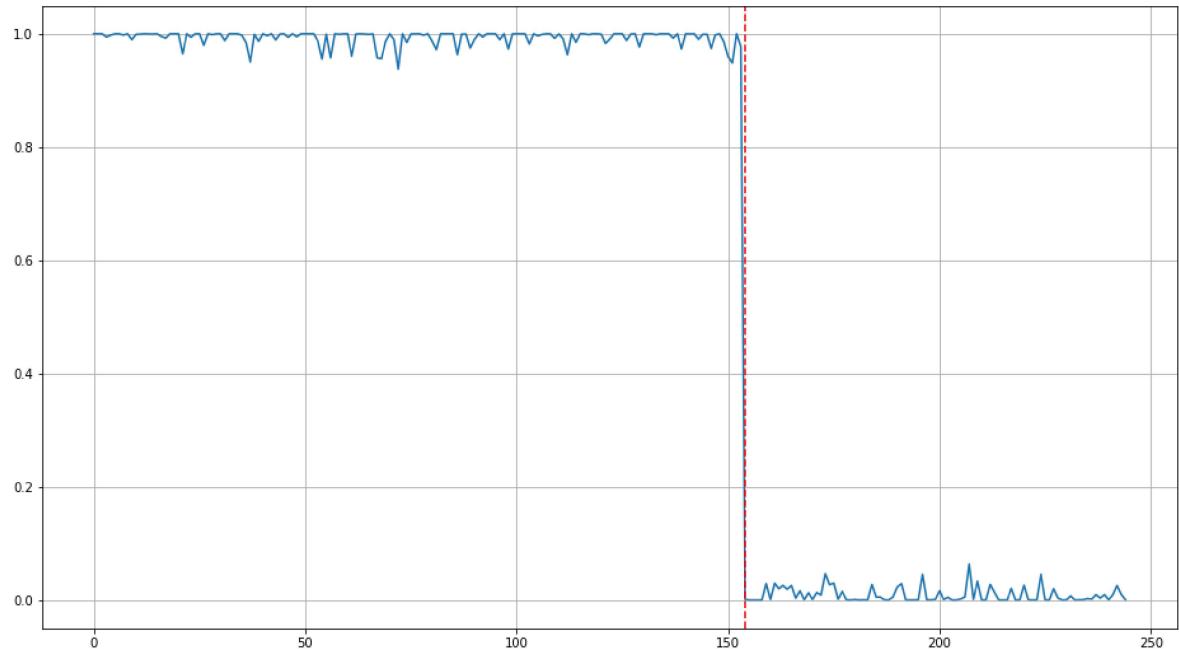
ax= plt.subplot()
sns.heatmap(cm, annot=True, fmt='g', ax=ax); #annot=True to annotate cells, f

# Labels, title and ticks
ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels');
ax.set_title('Confusion Matrix');
ax.xaxis.set_ticklabels(['Tumor','Healthy'])
ax.yaxis.set_ticklabels(['Tumor','Healthy'])
```

```
Out[72]: [Text(0, 0.5, 'Tumor'), Text(0, 1.5, 'Healthy')]
```



```
In [73]: plt.figure(figsize=(16,9))
plt.plot(outputs)
plt.axvline(x=len(tumor), color='r', linestyle='--')
plt.grid()
```



Visualising the Feature Maps of the Convolutional Filters

```
In [74]: model
```

```
Out[74]: CNN(
    (cnn_model): Sequential(
        (0): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
        (1): Tanh()
        (2): AvgPool2d(kernel_size=2, stride=5, padding=0)
        (3): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
        (4): Tanh()
        (5): AvgPool2d(kernel_size=2, stride=5, padding=0)
    )
    (fc_model): Sequential(
        (0): Linear(in_features=256, out_features=120, bias=True)
        (1): Tanh()
        (2): Linear(in_features=120, out_features=84, bias=True)
        (3): Tanh()
        (4): Linear(in_features=84, out_features=1, bias=True)
    )
)
```

```
In [75]: no_of_layers = 0
conv_layers = []
```

```
In [76]: model_children = list(model.children())
model_children
```

```
Out[76]: [Sequential(
    (0): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
    (1): Tanh()
    (2): AvgPool2d(kernel_size=2, stride=5, padding=0)
    (3): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
    (4): Tanh()
    (5): AvgPool2d(kernel_size=2, stride=5, padding=0)
), Sequential(
    (0): Linear(in_features=256, out_features=120, bias=True)
    (1): Tanh()
    (2): Linear(in_features=120, out_features=84, bias=True)
    (3): Tanh()
    (4): Linear(in_features=84, out_features=1, bias=True)
)]
```

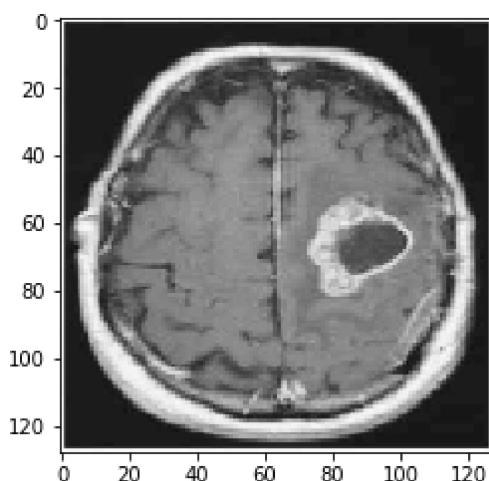
```
In [77]: for child in model_children:
    if type(child) == nn.Sequential:
        for layer in child.children():
            if type(layer) == nn.Conv2d:
                no_of_layers += 1
                conv_layers.append(layer)
```

```
In [78]: conv_layers
```

```
Out[78]: [Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1)),
Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))]
```

```
In [79]: img = mri_dataset[100]['image']
plt.imshow(img.reshape(128,128,3))
```

```
Out[79]: <matplotlib.image.AxesImage at 0x1d6cd612320>
```



```
In [80]: img = torch.from_numpy(img).to(device)
```

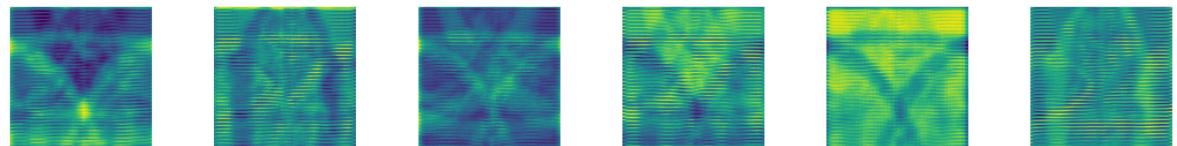
```
In [81]: img.shape
```

```
Out[81]: torch.Size([3, 128, 128])
```

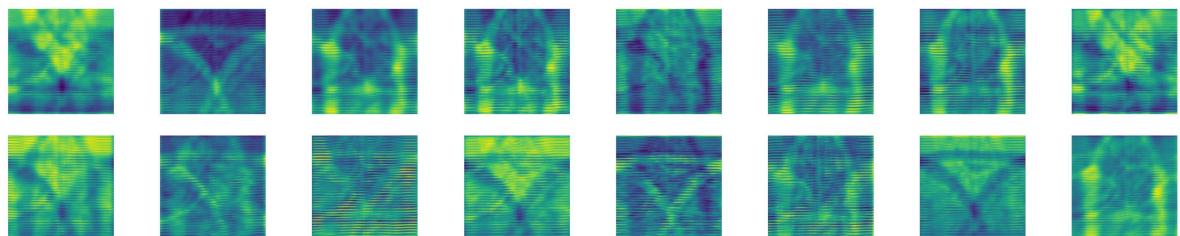
```
In [83]: img = img.unsqueeze(0) # Add batch dimension
results = [conv_layers[0](img)]
for i in range(1, len(conv_layers)):
    results.append(conv_layers[i](results[-1]))
outputs = results
```

```
In [84]: for num_layer in range(len(outputs)):
    plt.figure(figsize=(50, 10))
    layer_viz = outputs[num_layer].squeeze()
    print("Layer ", num_layer+1)
    for i, f in enumerate(layer_viz):
        plt.subplot(2, 8, i + 1)
        plt.imshow(f.detach().cpu().numpy())
        plt.axis("off")
    plt.show()
    plt.close()
```

Layer 1



Layer 2



Are We Over-fitting?

Preparing a validation set: We need to change the MRI dataset slightly!

We will need to make changes to our **MRI dataset class**:

- Define a function to divide the data into train and validation sets
- Define a variable called **mode** to determine whether we are interested in the training OR validation data
- Change **len()** and **getitem()** functions and conditioned over the variable **mode**

```
In [86]: # Import train/test split function from sklearn
from sklearn.model_selection import train_test_split
```


In [87]: `class MRI(Dataset):`

```

    def __init__(self):

        # Variables to hold the Training data and Validation data
        self.X_train, self.y_train, self.X_val, self.y_val = None, None, None,
        
        # A variable to determine if we are interested in retrieving the train
        self.mode = 'train'

        tumor = []
        healthy = []
        # cv2 - It reads in BGR format by default
        for f in glob.iglob("./data/brain_tumor_dataset/yes/*.jpg"):
            img = cv2.imread(f)
            img = cv2.resize(img,(128,128)) # I can add this later in the boot
            b, g, r = cv2.split(img)
            img = cv2.merge([r,g,b])
            img = img.reshape((img.shape[2],img.shape[0],img.shape[1])) # other
            tumor.append(img)

        for f in glob.iglob("./data/brain_tumor_dataset/no/*.jpg"):
            img = cv2.imread(f)
            img = cv2.resize(img,(128,128))
            b, g, r = cv2.split(img)
            img = cv2.merge([r,g,b])
            img = img.reshape((img.shape[2],img.shape[0],img.shape[1]))
            healthy.append(img)

        # our images
        tumor = np.array(tumor,dtype=np.float32)
        healthy = np.array(healthy,dtype=np.float32)

        # our labels
        tumor_label = np.ones(tumor.shape[0], dtype=np.float32)
        healthy_label = np.zeros(healthy.shape[0], dtype=np.float32)

        # Concatenates
        self.images = np.concatenate((tumor, healthy), axis=0)
        self.labels = np.concatenate((tumor_label, healthy_label))

    # Define a function that would separate the data into Training and Validation
    def train_val_split(self):
        self.X_train, self.X_val, self.y_train, self.y_val = \
        train_test_split(self.images, self.labels, test_size=0.20, random_state=42)

    def __len__(self):
        # Use self.mode to determine whether train or val data is of interest
        if self.mode == 'train':
            return self.X_train.shape[0]
        elif self.mode == 'val':
            return self.X_val.shape[0]

    def __getitem__(self, idx):
        # Use self.mode to determine whether train or val data is of interest
        if self.mode== 'train':
            sample = {'image': self.X_train[idx], 'label': self.y_train[idx]}

        elif self.mode== 'val':
            sample = {'image': self.X_val[idx], 'label': self.y_val[idx]}

```

```
    return sample

def normalize(self):
    self.images = self.images/255.0
```

Are we overfitting?

```
In [88]: mri_dataset = MRI()
mri_dataset.normalize()
mri_dataset.train_val_split()
```

```
In [89]: train_dataloader = DataLoader(mri_dataset, batch_size=32, shuffle=True)
val_dataloader = DataLoader(mri_dataset, batch_size=32, shuffle=False)
```

```
In [91]: device = torch.device("cpu")
model = CNN().to(device)
```

```
In [92]: eta=0.0001
optimizer = torch.optim.Adam(model.parameters(), lr=eta)
```

```
In [93]: # keep track of epoch losses
epoch_train_loss = []
epoch_val_loss = []
```

```
In [94]: for epoch in range(1,600):
    train_losses = []
    # train for the current epoch
    model.train()
    mri_dataset.mode = 'train'
    for D in train_dataloader:
        # Train the model
        optimizer.zero_grad()
        data = D['image'].to(device)
        label = D['label'].to(device)

        y_hat = model(data)
        error = nn.BCELoss()
        loss = torch.sum(error(y_hat.squeeze(), label))
        loss.backward()
        optimizer.step()
        train_losses.append(loss.item())

    epoch_train_loss.append(np.mean(train_losses))

    # validate for the current epoch
    val_losses = []
    model.eval()

    mri_dataset.mode = 'val'

    with torch.no_grad():
        for D in val_dataloader:
            data = D['image'].to(device)
            label = D['label'].to(device)
            y_hat = model(data)
            error = nn.BCELoss()
            loss = torch.sum(error(y_hat.squeeze(), label))
            val_losses.append(loss.item())

    epoch_val_loss.append(np.mean(val_losses))

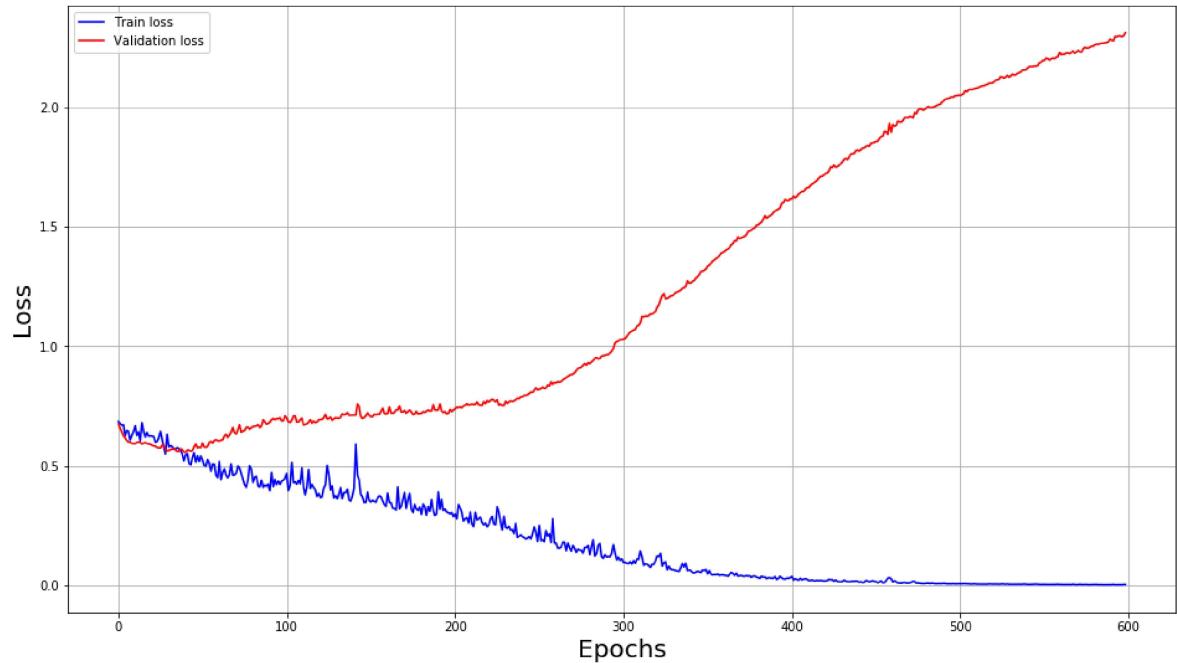
    if (epoch+1) % 10 == 0:
        print('Train Epoch: {} \tTrain Loss: {:.6f} \tVal Loss: {:.6f}'.format(e
```



Train Epoch: 10	Train Loss: 0.626618	Val Loss: 0.594669
Train Epoch: 20	Train Loss: 0.624329	Val Loss: 0.593146
Train Epoch: 30	Train Loss: 0.549134	Val Loss: 0.570088
Train Epoch: 40	Train Loss: 0.547733	Val Loss: 0.568139
Train Epoch: 50	Train Loss: 0.518123	Val Loss: 0.578808
Train Epoch: 60	Train Loss: 0.460436	Val Loss: 0.608882
Train Epoch: 70	Train Loss: 0.460993	Val Loss: 0.660111
Train Epoch: 80	Train Loss: 0.500268	Val Loss: 0.653505
Train Epoch: 90	Train Loss: 0.415764	Val Loss: 0.673874
Train Epoch: 100	Train Loss: 0.440570	Val Loss: 0.688613
Train Epoch: 110	Train Loss: 0.425302	Val Loss: 0.693222
Train Epoch: 120	Train Loss: 0.372940	Val Loss: 0.683721
Train Epoch: 130	Train Loss: 0.366797	Val Loss: 0.696388
Train Epoch: 140	Train Loss: 0.352535	Val Loss: 0.711878
Train Epoch: 150	Train Loss: 0.390249	Val Loss: 0.721253
Train Epoch: 160	Train Loss: 0.338560	Val Loss: 0.718994
Train Epoch: 170	Train Loss: 0.331603	Val Loss: 0.725447
Train Epoch: 180	Train Loss: 0.313739	Val Loss: 0.716998
Train Epoch: 190	Train Loss: 0.292859	Val Loss: 0.733707
Train Epoch: 200	Train Loss: 0.305369	Val Loss: 0.726395
Train Epoch: 210	Train Loss: 0.261251	Val Loss: 0.754153
Train Epoch: 220	Train Loss: 0.260625	Val Loss: 0.760360
Train Epoch: 230	Train Loss: 0.232124	Val Loss: 0.752555
Train Epoch: 240	Train Loss: 0.205698	Val Loss: 0.780609
Train Epoch: 250	Train Loss: 0.224908	Val Loss: 0.825945
Train Epoch: 260	Train Loss: 0.279764	Val Loss: 0.840991
Train Epoch: 270	Train Loss: 0.157038	Val Loss: 0.879134
Train Epoch: 280	Train Loss: 0.158401	Val Loss: 0.921043
Train Epoch: 290	Train Loss: 0.113376	Val Loss: 0.959788
Train Epoch: 300	Train Loss: 0.105093	Val Loss: 1.027755
Train Epoch: 310	Train Loss: 0.110376	Val Loss: 1.084217
Train Epoch: 320	Train Loss: 0.088628	Val Loss: 1.139928
Train Epoch: 330	Train Loss: 0.068093	Val Loss: 1.210473
Train Epoch: 340	Train Loss: 0.062575	Val Loss: 1.273066
Train Epoch: 350	Train Loss: 0.062688	Val Loss: 1.317345
Train Epoch: 360	Train Loss: 0.044383	Val Loss: 1.388746
Train Epoch: 370	Train Loss: 0.039089	Val Loss: 1.457390
Train Epoch: 380	Train Loss: 0.039524	Val Loss: 1.495853
Train Epoch: 390	Train Loss: 0.029998	Val Loss: 1.549515
Train Epoch: 400	Train Loss: 0.029658	Val Loss: 1.608582
Train Epoch: 410	Train Loss: 0.028185	Val Loss: 1.655849
Train Epoch: 420	Train Loss: 0.019675	Val Loss: 1.710487
Train Epoch: 430	Train Loss: 0.015302	Val Loss: 1.757073
Train Epoch: 440	Train Loss: 0.018502	Val Loss: 1.817910
Train Epoch: 450	Train Loss: 0.019577	Val Loss: 1.850106
Train Epoch: 460	Train Loss: 0.032678	Val Loss: 1.931545
Train Epoch: 470	Train Loss: 0.010090	Val Loss: 1.956541
Train Epoch: 480	Train Loss: 0.007997	Val Loss: 1.986927
Train Epoch: 490	Train Loss: 0.007958	Val Loss: 2.011108
Train Epoch: 500	Train Loss: 0.006886	Val Loss: 2.047710
Train Epoch: 510	Train Loss: 0.006984	Val Loss: 2.075544
Train Epoch: 520	Train Loss: 0.006132	Val Loss: 2.100352
Train Epoch: 530	Train Loss: 0.006213	Val Loss: 2.132497
Train Epoch: 540	Train Loss: 0.006059	Val Loss: 2.155070
Train Epoch: 550	Train Loss: 0.004082	Val Loss: 2.187640
Train Epoch: 560	Train Loss: 0.004882	Val Loss: 2.210636
Train Epoch: 570	Train Loss: 0.003554	Val Loss: 2.225332
Train Epoch: 580	Train Loss: 0.003674	Val Loss: 2.253994
Train Epoch: 590	Train Loss: 0.003566	Val Loss: 2.273471
Train Epoch: 600	Train Loss: 0.003930	Val Loss: 2.310722

```
In [95]: plt.figure(figsize=(16,9))
plt.plot(epoch_train_loss, c='b', label='Train loss')
plt.plot(epoch_val_loss, c='r', label = 'Validation loss')
plt.legend()
plt.grid()
plt.xlabel('Epochs', fontsize=20)
plt.ylabel('Loss', fontsize=20)
```

Out[95]: Text(0, 0.5, 'Loss')



In []: